

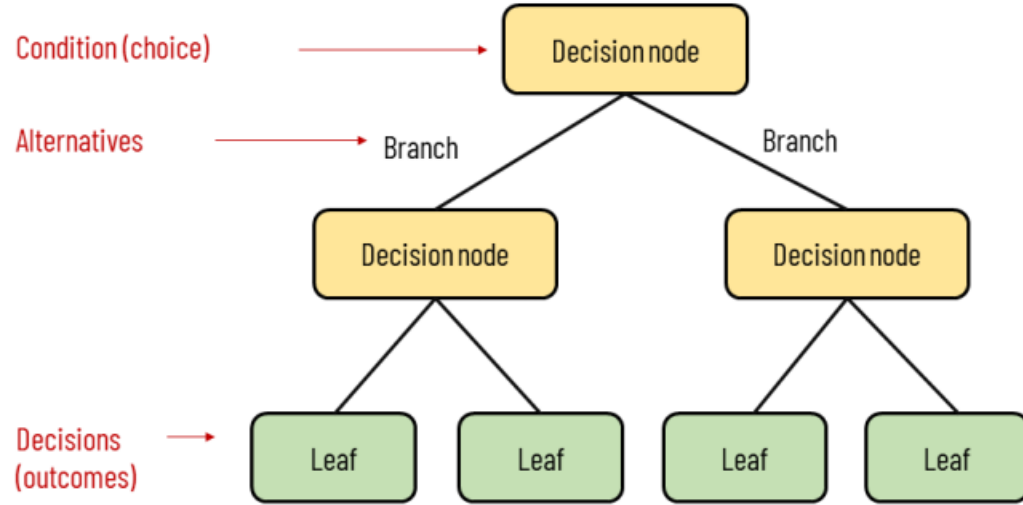


Random Forest

(Decision Trees)

NATHAN VAN SCHYNDEL, HAYDEN MUSCHA, AZURE ELLER, VICTOR PHAM

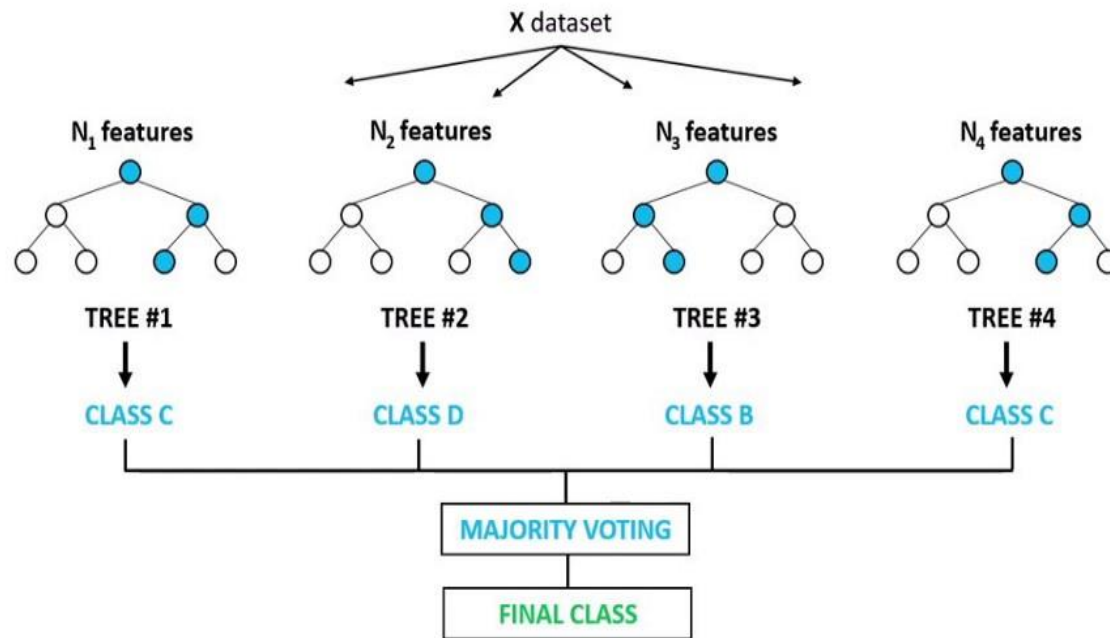
Elements of a decision tree



Random Forest

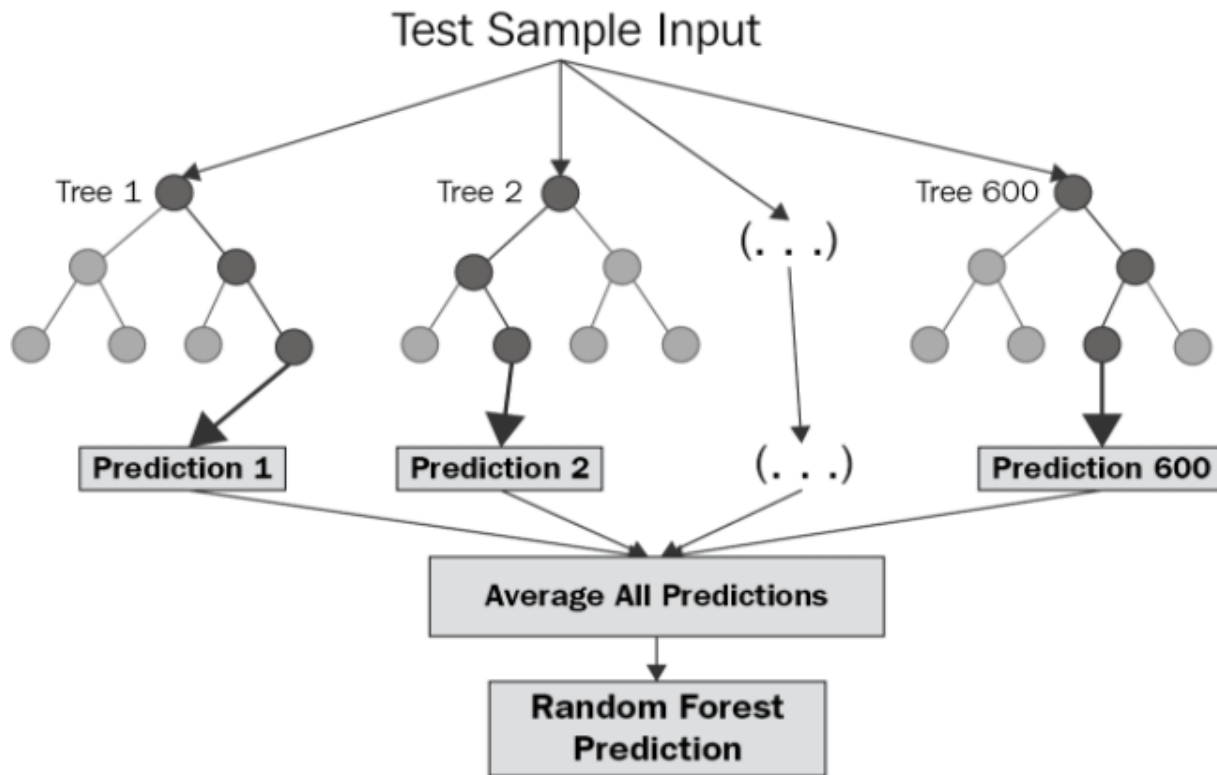
- Root Node
- Splitting
- Leaf
- Branch
- Tree

Random Forest Classifier



Forest Structure

- Number of trees
- Depth
- Voting Principle
 - Gini Impurity



Random Forest

- Statistical decisions

Forest:

Class Comparison

- Supervised
- Non-Linear or Linear
- Categorical and Continuous Specialty

Random Forest: Advantages

- Reduces Overfitting
- Flexible
- Categorical and Continuous
- Fills missing values
- No Normalization

Random Forest: Disadvantages

- Computationally heavy
- Time intensive
- Interpretability and variable significance

Data processing steps

- ▶ No Standardization
- ▶ Deal with null values, create dummy variables if necessary
- ▶ Split data from target variable
- ▶ Check feature importance



Hyperparameters: Regression

n_estimators - int, default=100

criterion - squared error, absolute error, poisson - default = squared error

max_depth - int, default=None

min_samples_split - int or float, default=2

min_samples_leaf - int or float, default=1

min_weight_fraction_leaf - float, default=0.0

max_features - {"sqrt", "log2", None}, int or float, default=1.0

max_leaf_nodes - int, default=None

min_impurity_decrease - float, default=0.0

bootstrap - bool, default=True

Hyperparameters: Classification

n_estimators - int, default=100

criterion - gini, entropy, log_loss, default = gini

max_depth- int, default=None

min_samples_split - int or float, default=2

min_samples_leaf - int or float, default=1

min_weight_fraction_leaf - float, default=0.0

max_features - {"sqrt", "log2", None}, int or float, default=sqrt

max_leaf_nodes - int, default=None

min_impurity_decrease - float, default=0.0

bootstrap - bool, default=True

Implementation of Random Forest Classifier ¶

Import the necessary libraries

```
In [79]: 1 import pandas as pd
2 from sklearn.metrics import confusion_matrix
3 from sklearn.metrics import matthews_corrcoef
4 # metrics are used to find accuracy or error
5 from sklearn import metrics
```

Load in the diabetes dataset

```
In [ ]: 1 data = pd.read_csv('diabetes.csv')
```

Split target feature from the rest of the data and run train_test_split

```
In [32]: 1 X = data.copy().drop(columns='Outcome')
2 y = data['Outcome'].copy()
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Set hyper parameters, fit the model on the training data, and then run predictions on the test data

```
In [108]: 1 clf = RandomForestClassifier(n_estimators = 50, max_features= 'log2', bootstrap = True, oob_score=0.78125)
2
3 # Training the model on the training dataset
4 # fit function is used to train the model using the training sets as parameters
5 clf.fit(X_train, y_train)
6
7 # performing predictions on the test dataset
8 y_pred = clf.predict(X_test)
```

Check R-squared

```
In [109]: 1 # using metrics module for accuracy calculation
2 print("Accuracy Score", clf.score(X_test,y_test))
```

Accuracy Score 0.78125

Sample Code:
Random Forest
Classifier on
Diabetes Dataset

Regression tree

Importing libraries

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
```

Read the Cars93 csv

```
In [ ]: cars = pd.read_csv('Cars93.csv')
```

Clean the cars data and get dummies

```
In [ ]: rowcleancars = cars.dropna()
ccars = rowcleancars.drop(columns=['Unnamed: 0'])
ccars = pd.get_dummies(ccars, drop_first=True)
ccars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 82 entries, 0 to 92
Columns: 225 entries, Min.Price to Make_Volvo 850
dtypes: float64(7), int64(11), uint8(207)
memory usage: 28.7 KB
```

Set up the testing and training data

```
In [ ]: X = ccars.drop(columns=['MPG.highway'])
y = ccars['MPG.highway']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Set up and train the regressor

```
In [ ]: regressor = RandomForestRegressor(n_estimators = 100, random_state = 0)
regressor.fit(X_train, y_train)
```

```
Out[ ]: RandomForestRegressor(random_state=0)
```

Performing prediction and checking r squared

```
In [ ]: pred = regressor.predict(X_test)
print("R-squared: ", regressor.score(X_test, y_test))
```

```
R-squared: 0.8179437894454142
```

Sample Code: Random Forest Regression Tree on Cars93

Summary

- ▶ An advancement of decision trees
- ▶ Can be used for both Regression and Classification
- ▶ Regression and Classification run with different criterion
- ▶ Easy to use
- ▶ Can be prone to overfitting if the hyperparameters are not managed

